

AI and Machine Learning Hardware Guide

AI workloads rely heavily on matrix multiplication performance.

CPUs handle orchestration while GPUs or NPUs accelerate computation.

Modern laptops include AI accelerators such as Intel NPU or AMD XDNA.

Memory bandwidth is often the bottleneck for inference.

Quantized models significantly reduce hardware requirements.